# Guidewise

Brand rules that run. Guidewise checks marketing content against a team's own approved rules and dated facts before it is published. Every finding quotes the rule it breaks and the exact words, every fix is re-tested before it is offered, risky content goes to the right reviewer, and a person makes the final approval.

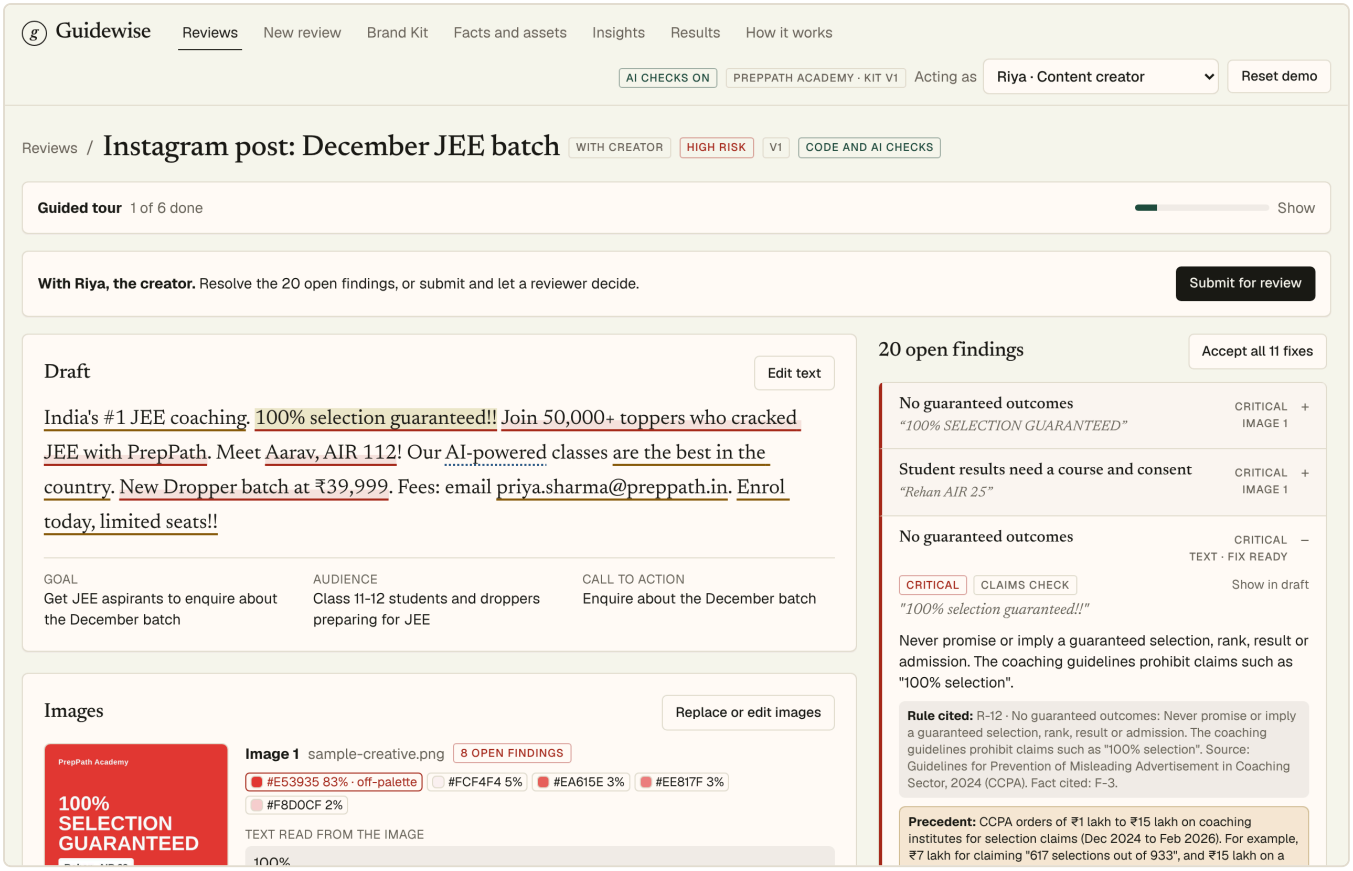
| | |
|---|---|
| Live product | **https://guidewise-kohl.vercel.app** · no sign-in; each visitor gets a private demo workspace |
| Walkthrough video | https://guidewise-kohl.vercel.app/demo.webm · 90 seconds, captioned, recorded on the live site |
| Source code | https://github.com/mahak867/Guide-wise |
| Demo brand | PrepPath Academy, a fictional JEE/NEET coaching brand written by the team |
| Team | Mahak Fahad (team leader, Backend + AI/ML) · Roshini R (Frontend + Design/UI-UX) · Chennai Institute of Technology |

## Where Round 2 stands

The central workflow runs end to end.
Paste copy or upload a DOCX, PDF or images → cited findings with re-tested fixes → accept, edit or dismiss with a reason → a fixed rule table routes to the brand or compliance reviewer → approval with a fingerprinted record. Tested end to end on the live site with the AI checks on.

Measured on our own test set.
On 40 held-out drafts with planted errors, code checks alone reach 100% precision and 74% recall; with the AI checks, 92% precision and 98% recall. 50 automated tests and CI run on every push.

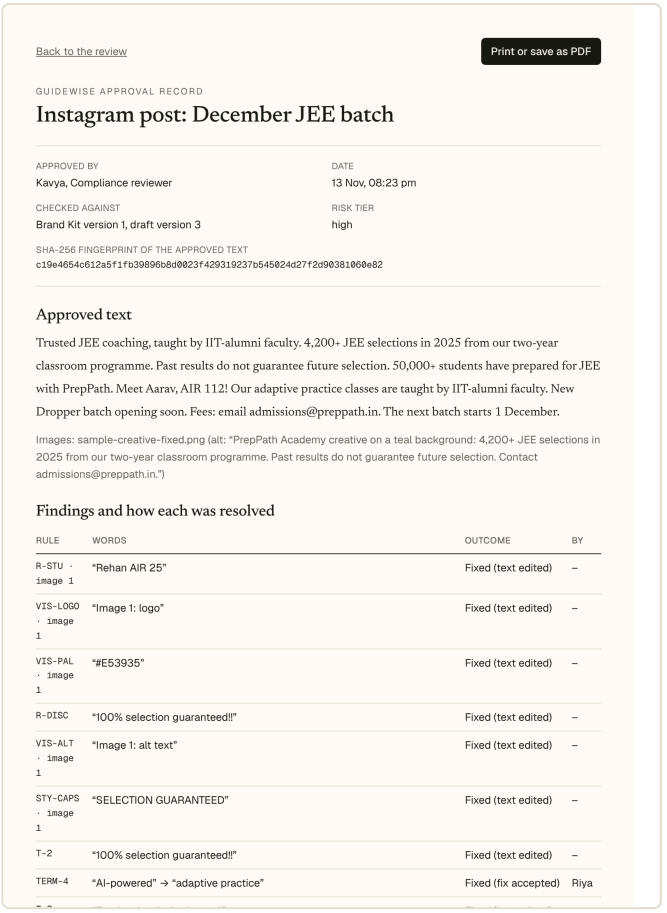
# The product



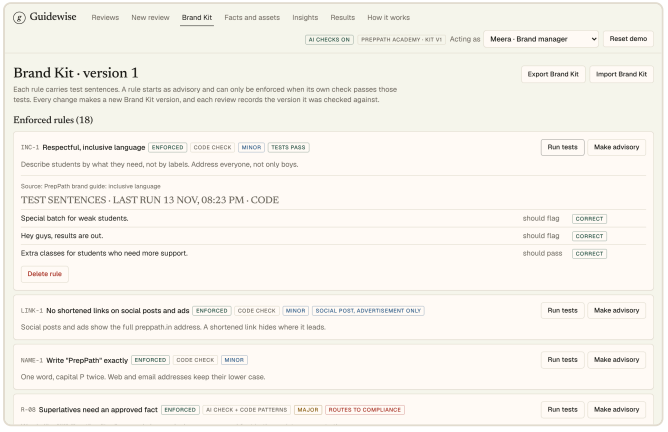

The sample Instagram post: the draft set as a manuscript with findings underlined, and one finding open with its rule, source, precedent and re-tested fix.



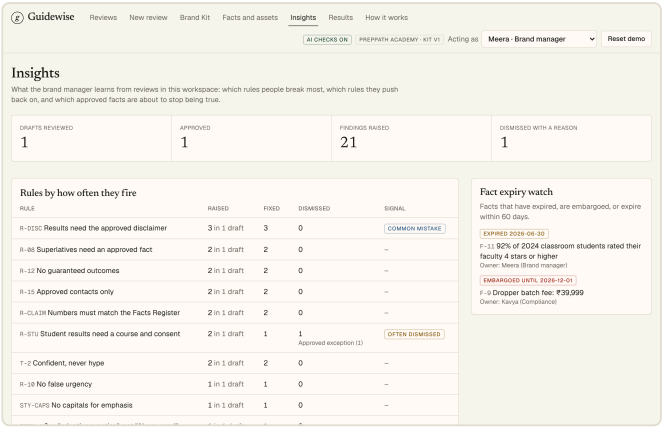

Text read from the creative is checked against every rule and fact; colours are checked against the palette.




Approval: the SHA-256 fingerprint of the exact approved text, the Brand Kit version, and every finding with how it was resolved.




Brand Kit: each rule carries test sentences and can only be enforced when its own check passes them.




Insights: which rules fire most, which are dismissed and why, and facts about to expire.

# Completed features

| AREA | WHAT A JUDGE CAN DO WITH IT |
|---|---|
| Review | Paste copy, or upload DOCX, PDF and up to four images; state goal, audience, call to action and locked phrases no fix may remove. |

| | |
|---|---|
| Code checks | Preferred terms, banned phrases, approved contacts, student results (course and consent record), the results disclaimer (with a same-language check for Hindi claims), and every figure against the Facts Register. |
| AI checks | Claims and voice rules judged one rule per call, run twice when they find something; disagreement becomes a suggestion. One adapter tries Claude, Groq and Gemini in that order, whichever keys are set (the live site runs on Groq, with Gemini as backup), with a 20-second cap per call; the app keeps working in rules-only mode if none answers. |
| Verifier | Plain code drops any finding whose quoted words, rule or fact does not exist ("no source, no flag"). |
| Images | OCR in the browser (served from our own site), palette check in CIELAB, WCAG contrast, alt text; image text runs through every rule and fact. Logo check: the Brand Kit logo is searched for at many sizes (masked normalised cross-correlation, no AI), then checked for minimum size and clear space. |
| Writing style | Code rule types for exact brand-name spelling, capitals used for emphasis, maximum length per content type (ad text cut off at "See more"), maximum sentence length, and inclusive wording, each with a fix where one is safe. |
| Facts Register | Facts with evidence, owner, valid-from, valid-until and required disclosure; embargoed, expired and wrong-subject figures are each flagged. |
| Fixes | Each fix is applied and re-checked before it is offered; AI-written fixes may not introduce any violation; edit a fix before accepting; "Accept all" works through follow-ups. |
| Brand Kit | Rules with test sentences, an enforce gate, versioning, and a guide compiler that drops any proposed rule whose quoted passage is not in the guide. |
| Regulation pack | Central Consumer Protection Authority (CCPA, India) guidelines for the coaching sector (2024), with a penalty order shown on each regulation finding as reported in the press, and the "more than 50 students" applicability gate. |
| Workflow | Four roles enforced on the server; risk routing by a fixed rule table; comments, version history with word-level diffs, activity log, approval record (save as PDF). |
| Insights | Most-flagged rules, dismissal reasons, rules two or more people call unclear, fact expiry watch. |
| Privacy | Emails, phone numbers, Aadhaar, PAN and card numbers are masked in code before any AI call. A draft marked confidential is never sent to AI at all. |
| Channels | A rule can apply to some content types only (for example, no shortened links on social posts and ads); the brief card shows the rules for the type being written. |
| Portability | The Brand Kit exports and imports as one JSON file; imported new or changed rules arrive as advisory until they pass their tests. |
| Reliability | AI answers are cached by rule and text, so many judges reviewing the same sample spend no extra quota; a 20-second cap on every AI call; a 20-step browser test of the whole journey runs against the live site. |
| Usability | One-click sample post, a guided tour that ticks itself off, private workspace per visitor, works on a phone. |

## How it is built

| 1 Redact | 2 Code checks | 3 AI checks | 4 Verifier | 5 Re-test fixes | 6 Route |
|---|---|---|---|---|---|
| Personal data out before any AI call | Terms, contacts, students, every figure vs the register | One rule per call, run twice | Quote, rule and fact must exist | New issue → fix withdrawn | Fixed rule table, human approves |

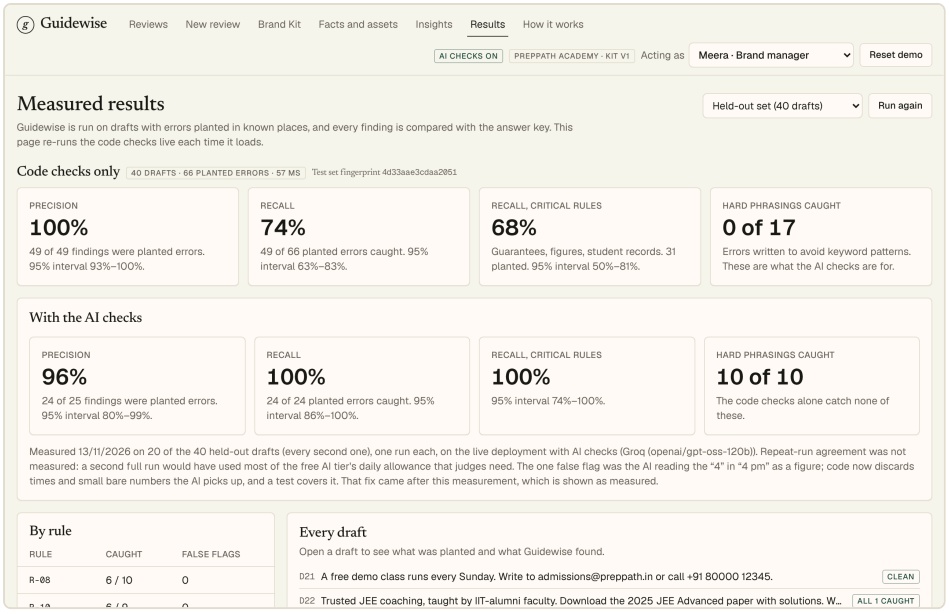
Only step 3 uses AI. The AI only locates number claims; code re-reads each value and compares it with the Facts Register.

| LAYER | CHOICE AND REASON |
|---|---|
| App and API | Next.js 16 (TypeScript), React 19, Tailwind CSS 4 on Vercel, functions in Mumbai next to the database. One codebase, one language for a team of two. |
| Data | Supabase Postgres; every table scoped by a workspace ID so each visitor's demo is isolated. |
| AI | One adapter for Groq (openai/gpt-oss-120b), Gemini and Claude with automatic fallback; every answer must pass a Zod schema. |
| Documents and images | mammoth (DOCX), unpdf (PDF), tesseract.js (OCR) — read in memory, never stored. |
| Quality | 50 Vitest tests, including eight that swap in a deliberately wrong AI (invented quotes, a forbidden fix, a wrong number, a time misread as a figure, personal data, a confidential draft, prompt injection, locked phrases); CI runs lint, type-check, tests, evaluation and build on every push. |

## Measured results

60 short drafts for PrepPath with errors planted in known places (20 used while building, **40 held out**). A finding counts as correct when it cites the planted error's rule and quotes overlapping words; anything else is a false flag. Intervals are 95% Wilson intervals. Test-set fingerprint `4d33aae3cdaa2051`.

| HELD-OUT SET | PRECISION | RECALL | CRITICAL RECALL | HARD PHRASINGS | REPEAT-RUN AGREEMENT |
|---|---|---|---|---|---|
| Code checks only (40 drafts) | 100% (93%–100%) | 74% (63%–83%) | 68% | 0 of 17 | n/a (no randomness) |
| Code and AI checks (40 drafts) | 92% (83%–96%) | 98% (92%–100%) | 100% | 17 of 17 | 83% |




The live Results page: the code checks re-run on every visit; the AI checks were measured on the live site.

The AI row is measured on all 40 held-out drafts on the live site, run twice with the answer cache off; the first run is reported, and the two runs agree on 83% of drafts. Of its false flags, four are the AI calling "cheapest" or "Best" an unsupported superlative (R-08) where the answer key planted only the "cheap" term rule; we count them as false rather than change the key after seeing the results. One flagged "Last chance" under the tone rule instead of the urgency rule, which is also its one miss. One read a rating ("4 stars") as a figure; code now discards rating-scale values the AI picks up, with a test, but that fix came after this run, which is shown as measured.

These numbers are optimistic: the team wrote both the rules and the test drafts. "Hard" errors were written on purpose to avoid keyword patterns. The next step is real, anonymised drafts from coaching-centre marketers.

## Changes from the Round 1 plan

| PLAN | WHAT WE DID | WHY |
|---|---|---|
| Supabase Auth with four roles | No sign-in; a private workspace per visitor and a role switcher, role rules enforced on the server; browsers never receive database credentials | Judges must be able to open the link directly. Sign-in and per-organisation database security come with a production deployment. |
| Image review in Round 3 | A first version ships now (palette, contrast, alt text, logo presence, size and clear space, OCR text through every rule) | The brief names visual identity as part of the central workflow. Round 3 deepens it. |
| AI on Claude | One adapter: Claude, Groq, Gemini in that order, whichever keys are set; the live site runs on Groq with Gemini as backup | Free tiers make the AI checks available to judges at no cost; the schema and verifier are the same for all. |
| Fixes withdrawn on any new issue | Withdrawn only when the new issue has no automatic fix (AI-written fixes: any new issue) | Fixing a guarantee correctly raises a disclaimer requirement; withdrawing that fix was wrong. |
| alex, write-good, axe-core | Replaced by our own code rules: brand-name spelling, capitals, length, sentence length, inclusive wording; contrast and alt text for images | As Brand Kit rules they carry test sentences, can be edited by the manager and cite their source, which the libraries' fixed word lists could not. |
| Hands-on tests with students | Moved to the start of Round 3 | Build time went into image review and the measured evaluation. |

> **Correction to Round 1.** The Round 1 PDF stated CCPA maximum penalties of "up to ₹10 lakh / ₹50 lakh" on pages 2 and 14. We could not confirm these amounts against the Act and withdraw them. The penalty orders shown in the product are cited as reported in the linked press coverage.

## The three major improvements for Round 3

### 1. Guideline Health loop — the brand guide fixes itself

Round 2 already flags rules that two or more people dismiss as unclear. Round 3 closes the loop: Guidewise drafts an amendment with new test sentences, the manager approves it, and a new Brand Kit version must pass its tests. Changing a fact lists every draft and approved piece that used it.

Test: dismiss one rule as three different creators → it is marked unclear with a drafted amendment that passes its tests; change a fact → every affected piece is listed.

### 2. Governance at the source — an API and MCP server for AI writing tools

A REST API with scoped keys and an MCP server exposing `get_brief` and `check_draft`, so AI writing tools fetch the rules and the facts valid today and check a draft before handing it over. An in-app playground and a per-key audit trail.

Test: call `check_draft` from the playground and from an MCP client with a key → the same cited findings as the app, and a new audit entry.

### 3. Campaign and creative review with AI vision

Review a whole campaign together (caption, creatives, landing-page copy, PDF brochure) and flag inconsistencies across assets: prices, dates, claims, promo codes. AI vision checks what OCR cannot see: a student photo with no consent record, a recoloured or redrawn logo, stylised claims.

Test: upload a three-asset campaign with a price mismatch and an unconsented student photo → both found, cited and routed to compliance.

## Team and disclosures

| MEMBER | ROLE | COLLEGE | COURSE |
|---|---|---|---|
| Mahak Fahad (team leader) | Backend + AI / ML | Chennai Institute of Technology | B.E. Computer Science and Engineering, batch of 2029 |
| Roshini R | Frontend + Design / UI-UX | Chennai Institute of Technology | B.E. Computer Science and Engineering, batch of 2029 |

- **AI used in building:** Claude (Anthropic) was used to help write code and documents during the official build window. The team reviewed it and is responsible for it.
- **AI in the product:** Groq and Gemini APIs (Claude supported), through one adapter. Personal data is masked before any call.
- **Reused code or templates:** none, apart from the Next.js starter (create-next-app). All product code was written inside the build window; the commit history starts on 13 November 2026 at 10:09 IST.
- **Open source:** Next.js, React, Tailwind CSS, supabase-js, Zod, Anthropic SDK, groq-sdk, @google/genai, mammoth, unpdf, tesseract.js.
- **Demo data:** PrepPath Academy, its rules, facts and students are fictional. Penalty orders are cited as reported in the linked press coverage; we have not read the orders themselves. Guidewise supports reviewers; it does not give legal advice.